

Stochastic Simulation (MIE1613H) - Homework 3 (Solutions)

Due: March 21, 2022

Problem 1. (30 Pts.) You are asked to make projections about cycle times for a semiconductor manufacturer who plans to open a new plant. Here “cycle time” means the time from product release until completion. The process that you will consider has a single diffusion step with sub-steps as indicated in the diagram.

Raw material for two products (C) and (D) will begin at the CLEAN step, and make multiple passes until the product completes processing. Movement within each process is handled by robots and takes very little time (treat as 0) relative to the processing steps. The movement time from release to diffusion is 15 minutes (1/4 hour).

The anticipated product mix is to have 60% of products to be of type (C) and 40% of type (D). Product (C) requires 5 passes and product D requires 3 passes.

The OXDIZE step is deterministic but differs by product type: it takes 2.7 hours for (C) and 2.0 hours for (D).

(a) The CLEAN and LOAD/UNLOAD steps do not differ by product type but are subject to uncertainty. Historical data is provided in ([SemiconductorData.xls](#)). Represent them using the distributions available in PythonSim, and justify your choice using the graphical and statistical methods discussed in the lecture.

Product will be released in cassettes at the rate of 1 cassette/hour, 7 days a week, 24-hours a day (this is to achieve a desired throughput of 1 cassette/hour). Product is moved and processed in single cassette loads.

The table below shows the number of machines that are planned for each fabrication step:

Step	Number of Machines
CLEAN	9
LOAD QTZ	2
OXD	11
UNLOAD QTZ	2

Table 1: Number of available machines in each step.

Since we do not have a physical justification for the choice of distributions, we can speculate a few candidate distributions (e.g., by looking at the shape of the histogram) in each case and examine the goodness of fit for them. For the CLEAN data, we pick Exponential, Weibull, and Gamma distributions. The KS test does not reject any of the distributions at $(1 - \alpha) = 95\%$ confidence level. Examining the Q-Q plot we observe that the fits are very similar, so any of them would provide an acceptable choice. We pick exponential that is available in PythonSim. The estimated parameter for exponential is $\lambda = 0.5512835$, so the estimate of mean, that is used in PythonSim, is $1/0.5512835 = 1.813949$.

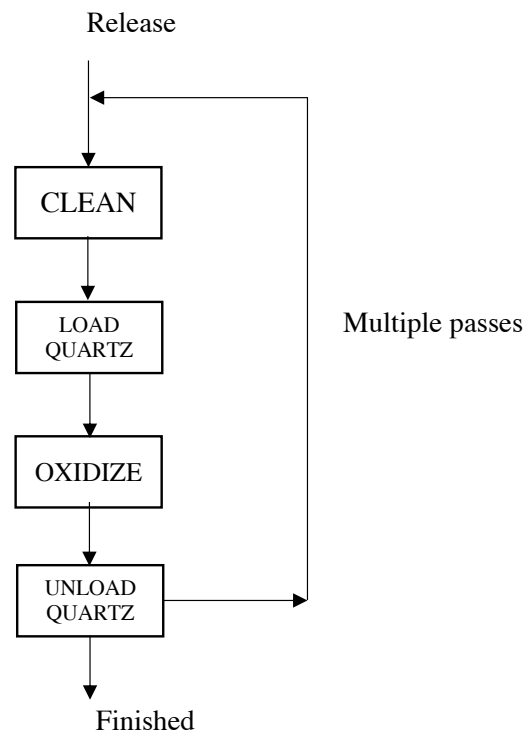


Figure 1: Diagram of the diffusion step.

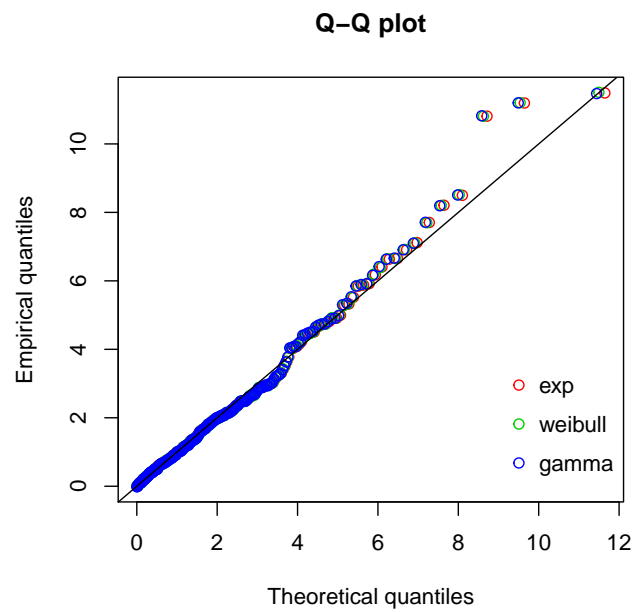


Figure 2: Q-Q plot for the fitted distributions to the CLEAN data.

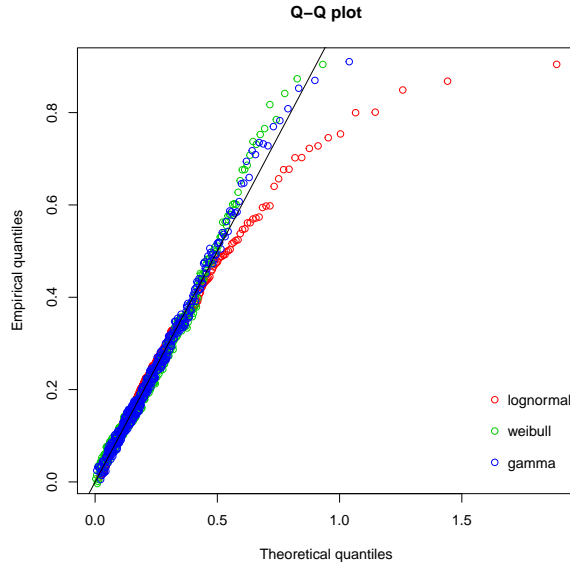


Figure 3: Q-Q plot for the fitted distributions to the LOAD data.

For the LOAD/UNLOAD data, we pick Lognormal, Weibull, and Gamma. The KS test does not reject any of the distributions. Examining the Q-Q plot however suggests that Gamma provides a better fit at the tail. However, we pick Lognormal that is available in PythonSim. Let X be the Lognormal random variable and denote by μ and σ the mean and standard deviation of $\log(X)$ as estimated in R. The mean and standard deviation of X , that is required in PythonSim, are given by,

$$E(X) = e^{\mu + \frac{\sigma^2}{2}} = 0.2475798,$$

$$SD(X) = e^{\mu + \frac{\sigma^2}{2}} \sqrt{e^{\sigma^2} - 1} = 0.2233498.$$

```

1  # This file uses the fitdistrplus package. You'll need to install it first
2  # by running the command: 'install.packages("fitdistrplus")' in R
3  require(fitdistrplus)
4
5  # import the data from .csv file
6  MyData <- read.csv("SemiconductorData.csv", header=FALSE)
7  Load <- MyData[,2]
8  Clean <- MyData[1:299,4]
9  plotdist(Load, histo = TRUE, demp = TRUE)
10 plotdist(Clean, histo = TRUE, demp = TRUE)
11
12 # fit a lognormal, gamma, and weibull to load data
13 load_fln <- fitdist(Load, "lnorm")
14 load_fw <- fitdist(Load, "weibull")
15 load_fg <- fitdist(Load, "gamma")
16
17 # fit an exponential, gamma, and weibull to clean data
18 clean_fe <- fitdist(Clean, "exp")

```

```

19 clean_fw <- fitdist(Clean, "weibull")
20 clean_fg <- fitdist(Clean, "gamma")
21
22 # compare the fits for load
23 plot.legend <- c("lognormal", "weibull", "gamma")
24 qqcomp(list(load_fln, load_fw, load_fg), legendtext = plot.legend)
25
26 # compare the fits for clean
27 plot.legend <- c("exp", "weibull", "gamma")
28 qqcomp(list(clean_fe, clean_fw, clean_fg), legendtext = plot.legend)
29
30 # Perform GofFit tests
31 gof_load <- gofstat(list(load_fln, load_fw, load_fg), fitnames = c("lognormal"
    , "weibull", "gamma"))
32 gof_load
33 gof_load$kstest
34
35 gof_clean <- gofstat(list(clean_fe, clean_fw, clean_fg), fitnames = c("exp", "
    weibull", "gamma"))
36 gof_clean
37 gof_clean$kstest

```

(b) Provide estimates of the long-run average cycle times for each product. Determine and justify a warm-up period and run length. You may use the replication-deletion approach. Provide appropriate confidence intervals for your estimates.

We define a special entity class in SimClasses.py as follows to represent a product.

```

1 # special entity to represent a product
2 class Product():
3     def __init__(self, type):
4         self.CreateTime = Clock
5         self.Type = type
6         if self.Type == "C":
7             self.Passes = 5
8             self.Oxd = 2.7
9         else:
10            self.Passes = 3
11            self.Oxd = 2.0

```

When a product is released, we schedule the next release for the next hour, determine the product type with the given probabilities, and schedule the end of movement from release to diffusion using the SchedulePlus function. In the EndMove function, we schedule the end of cleaning if a cleaning machine is available using the estimated exponential distribution from the data. Otherwise, we put the product in the cleaning queue. In the EndClean function, we schedule the end of loading if a loading machine is available using the estimated Lognormal distribution from the data. Otherwise, we put the product in the loading queue. We also check the cleaning queue and if there is a product waiting, we schedule the next end of cleaning for that product. Otherwise, we make a cleaning machine idle. The EndLoad and EndOxd functions are also similar to the EndClean function. In the EndUnload function, we keep track of the product passes and if a product reaches the required passes, we record the cycle time for that product type in its corresponding DTStat object. Otherwise, we check the status of the cleaning machines and repeat the same steps.




Figure 4: Mean plot, by product number, averaged across $n = 30$ replications.

Fig. 4 shows the mean cycle time by product number for the first $m = 1000$ products of each type averaged across $n = 30$ replications. At somewhere around 100 products the average cycle times seem to vary around a central value, so we may take $d = 100$ and choose the run length at least $m = 1100$ of each product. The estimate of the average cycle times for product C and D is 31.296 and 16.700 hours with the 95% CI of [31.282, 31.309] and [16.693, 16.708], respectively.

```

1  # Cycle times for a semiconductor manufacturer
2  import SimFunctions
3  import SimRNG
4  import SimClasses
5  import scipy.stats as sp
6  import numpy as np
7  import pandas as pd
8  import matplotlib.pyplot as plt
9
10 def mean_confidence_interval(data, confidence=0.95):
11     a = 1.0*np.array(data)
12     n = len(a)
13     m, se = np.mean(a), sp.sem(a)
14     h = se * sp.t._ppf((1+confidence)/2., n-1)
15     return m, m-h, m+h
16
17 ZSimRNG = SimRNG.InitializeRNSeed()
18 Calendar = SimClasses.EventCalendar()
19
20 # set up queues
21 CleanQ = SimClasses.FIFOQueue()
```

```

22 LoadQ = SimClasses.FIFOQueue()
23 OxdQ = SimClasses.FIFOQueue()
24 UnloadQ = SimClasses.FIFOQueue()
25 # set up resources
26 Clean = SimClasses.Resource()
27 Load = SimClasses.Resource()
28 Oxd = SimClasses.Resource()
29 Unload = SimClasses.Resource()
30
31 Clean.SetUnits(9)
32 Load.SetUnits(2)
33 Oxd.SetUnits(11)
34 Unload.SetUnits(2)
35
36 # set up statistics
37 CycleTimeC = SimClasses.DTStat()
38 CycleTimeD = SimClasses.DTStat()
39
40 TheCTStats = []
41 TheDTStats = [CycleTimeC, CycleTimeD]
42 TheQueues = [CleanQ, LoadQ, OxdQ, UnloadQ]
43 TheResources = [Clean, Load, Oxd, Unload]
44
45
46 def Release():
47     SimFunctions.Schedule(Calendar, "Release", 1.0)
48     if SimRNG.Uniform(0, 1, 1) < 0.6:
49         NewProduct = SimClasses.Product("C")
50     else:
51         NewProduct = SimClasses.Product("D")
52
53     SimFunctions.SchedulePlus(Calendar, "EndMove", 0.25, NewProduct)
54
55 def EndMove(TheProduct):
56     if Clean.Busy < Clean.NumberOfUnits:
57         Clean.Seize(1)
58         SimFunctions.SchedulePlus(Calendar, "EndClean", SimRNG.Expon(1.813949,
59             2), TheProduct)
60     else:
61         CleanQ.Add(TheProduct)
62
63 def EndClean(TheProduct):
64     if Load.Busy < Load.NumberOfUnits:
65         Load.Seize(1)
66         SimFunctions.SchedulePlus(Calendar, "EndLoad", SimRNG.Lognormal
67             (0.2475798, 0.2233498 ** 2, 3), TheProduct)
68     else:
69         LoadQ.Add(TheProduct)
70
71     if CleanQ.NumQueue() > 0:
72         NextProduct = CleanQ.Remove()
73         SimFunctions.SchedulePlus(Calendar, "EndClean", SimRNG.Expon(1.813949,
74             2), NextProduct)

```

```

73     else:
74         Clean.Free(1)
75
76
77 def EndLoad(TheProduct):
78     if Oxd.Busy < Oxd.NumberOfUnits:
79         Oxd.Seize(1)
80         SimFunctions.SchedulePlus(Calendar, "EndOxd", TheProduct.Oxd,
81                                     TheProduct)
82     else:
83         OxdQ.Add(TheProduct)
84
85     if LoadQ.NumQueue() > 0:
86         NextProduct = LoadQ.Remove()
87         SimFunctions.SchedulePlus(Calendar, "EndLoad", SimRNG.Lognormal
88                                     (0.2475798, 0.2233498 ** 2, 3), NextProduct)
89     else:
90         Load.Free(1)
91
92
93 def EndOxd(TheProduct):
94     if Unload.Busy < Unload.NumberOfUnits:
95         Unload.Seize(1)
96         SimFunctions.SchedulePlus(Calendar, "EndUnload", SimRNG.Lognormal
97                                     (0.2475798, 0.2233498 ** 2, 3), TheProduct)
98     else:
99         UnloadQ.Add(TheProduct)
100
101     if OxdQ.NumQueue() > 0:
102         NextProduct = OxdQ.Remove()
103         SimFunctions.SchedulePlus(Calendar, "EndOxd", NextProduct.Oxd,
104                                     NextProduct)
105     else:
106         Oxd.Free(1)
107
108
109 def EndUnload(TheProduct):
110     global CTC, CTD, TC, TD
111     TheProduct.Passes = TheProduct.Passes - 1
112     if TheProduct.Passes > 0:
113         if Clean.Busy < Clean.NumberOfUnits:
114             Clean.Seize(1)
115             SimFunctions.SchedulePlus(Calendar, "EndClean", SimRNG.Expon
116                                         (1.813949, 2), TheProduct)
117         else:
118             CleanQ.Add(TheProduct)
119     else:
120         if TheProduct.Type == "C":
121             CycleTimeC.Record(SimClasses.Clock - TheProduct.CreateTime)
122             CTC.append(SimClasses.Clock - TheProduct.CreateTime)
123             TC += 1
124         else:
125             CycleTimeD.Record(SimClasses.Clock - TheProduct.CreateTime)
126             CTD.append(SimClasses.Clock - TheProduct.CreateTime)

```

```

122         TD += 1
123
124     if UnloadQ.NumQueue() > 0:
125         NextProduct = UnloadQ.Remove()
126         SimFunctions.SchedulePlus(Calendar, "EndUnload", SimRNG.Lognormal
            (0.2475798, 0.2233498 ** 2, 3), NextProduct)
127     else:
128         Unload.Free(1)
129
130 AcrossC = []
131 AcrossD = []
132 AllCTC = [] # these will be a list of lists, each
133 AllCTD = [] # list corresponding to one replication's cycle times
134
135 m = 1100 #
136 d = 100 # Deletion point
137 for reps in range(0, 30, 1):
138     CTC = [] # cycle times from current replication
139     CTD = [] # cycle times from current replication
140     SimFunctions.SimFunctionsInit(Calendar, TheQueues, TheCTStats, TheDTStats,
        TheResources)
141     SimFunctions.Schedule(Calendar, "Release", 1.0)
142
143     NextEvent = Calendar.Remove()
144     SimClasses.Clock = NextEvent.EventTime
145     Release()
146
147     TC = 0 # Number of completed products of type C
148     TD = 0 # Number of completed products of type D
149     while min(TC, TD) < m:
150         if min(TC, TD) == d:
151             SimFunctions.ClearStats(TheCTStats, TheDTStats)
152
153         NextEvent = Calendar.Remove()
154         SimClasses.Clock = NextEvent.EventTime
155         if NextEvent.EventType == "Release":
156             Release()
157         elif NextEvent.EventType == "EndMove":
158             EndMove(NextEvent.WhichObject)
159         elif NextEvent.EventType == "EndClean":
160             EndClean(NextEvent.WhichObject)
161         elif NextEvent.EventType == "EndLoad":
162             EndLoad(NextEvent.WhichObject)
163         elif NextEvent.EventType == "EndOxd":
164             EndOxd(NextEvent.WhichObject)
165         elif NextEvent.EventType == "EndUnload":
166             EndUnload(NextEvent.WhichObject)
167
168
169     AcrossC.append(CycleTimeC.Mean())
170     AcrossD.append(CycleTimeD.Mean())
171     AllCTC.append(CTC)
172     AllCTD.append(CTD)
173

```



```

174 # output results
175 print("Estimated_expected_cycle_time_for_product_C:", mean_confidence_interval
      (AcrossC, 0.05))
176 print("Estimated_expected_cycle_time_for_product_D:", mean_confidence_interval
      (AcrossD, 0.05))
177
178 # Create mean plot to determine the warm-up period
179 MeanC = []
180 MeanD = []
181 for i in range(1000): # Average cycle time for the first 1000 products
182     MeanC.append(np.mean([rep[i] for rep in AllCTC]))
183     MeanD.append(np.mean([rep[i] for rep in AllCTD]))
184 plt.plot(MeanC, label = "Product_C")
185 plt.plot(MeanD, label = "Product_D")
186 plt.xlabel("Product")
187 plt.ylabel("Average_cycle_time")
188 plt.legend()
189 plt.show()

```

Estimated expected cycle time for product C: (31.29586, 31.28239, 31.30934)
Estimated expected cycle time for product D: (16.70039, 16.69263, 16.70815)

Problem 2. (20 Pts.) An Automated external defibrillator (AED) can save a person's life in event of a cardiac arrest. To accelerate delivery, start ups have been developing drone technology to quickly deliver AEDs to the scene of the cardiac arrest in case of an emergency call. The AEDs will be maintained at various bases across the city to respond to calls in the area designated to each base.

You have been tasked to estimate the minimum number of drones at a certain base to ensure that with 98% probability there will be a drone available at the event of a cardiac arrest in the area covered by that base.

Assume that calls reporting cardiac arrests arrive according to a non-homogeneous Poisson process with (per minute) rate function,

$$\lambda(t) = \begin{cases} 4, & 7\text{AM}-12\text{PM}, \\ 2, & 12\text{PM}-12\text{AM}, \\ 1, & 12\text{AM}-7\text{AM}. \end{cases}$$

throughout a day. Further, the total "drone busy time" i.e. from the time the call is received until the drone is back to the base and ready to dispatch again is exponentially distributed with mean 45 minutes. Assume at time 12AM, there are no calls (requests for AEDs) in the system.

(a) Explain, in words, your approach in determining the minimum number of drones at the base required to satisfy the provided service level. Specify the queueing model, corresponding parameters, and the performance measure you are estimating.

The approach is the same as the parking lot example and the $M(t)/M/\infty$ queue. The piecewise-constant arrival rate function $\lambda(t)$ can be used together with the thinning method to generate the NSPP. Note that the inter-arrival times are no longer exponentially distributed and hence generating exponential inter-arrival times with time-dependent rate is NOT a valid approach. The service times are i.i.d exponential with mean 45 minutes. The performance measure of interest is the 0.98th quantile of the maximum number of calls reporting cardiac arrests during the day.

(b) Determine the minimum number of drones at the base using the inputs provided. Provide a 95% CI for your estimate.

We simulate 1000 samples of the maximum number of calls during the day, sort them, and pick the 980th value as the minimum number of drones required at the base for the desired service level. The estimated required drones is 224 with the 95% CI of $[Y_{(971)}, Y_{(989)}] = [222, 226]$.

```

1  # Drones
2  import SimClasses
3  import SimFunctions
4  import SimRNG
5  import math
6  import pandas
7  import numpy as np
8
9  ZSimRNG = SimRNG.InitializeRNSeed()
10 Calendar = SimClasses.EventCalendar()
11 TheCTStats = []
12 TheDTStats = []
13 TheQueues = []
14 TheResources = []
15 AllMaxQueue = []
16
17 MeanBusyTime = 45.0
18 RepNum = 1000
19
20 ArrivalRates = [1, 4, 2]
21 MaxRate = 4
22
23 # Piecewise-constant arrival rate function
24 def PW_ArrRate(t):
25     if t < 7*60:      # 12AM-7AM
26         return ArrivalRates[0]
27     elif t < 12*60:   # 7AM-12PM
28         return ArrivalRates[1]
29     else:             # 12PM-12AM
30         return ArrivalRates[2]
31
32 # Thinning method used to generate NSPP with the piecewise-constant arrival
rate
33 def NSPP(Stream):
34     PossibleArrival = SimClasses.Clock + SimRNG.Expon(1/MaxRate, Stream)
35     while SimRNG.Uniform(0, 1, Stream) >= PW_ArrRate(PossibleArrival)/MaxRate:
36         PossibleArrival = PossibleArrival + SimRNG.Expon(1/MaxRate, Stream)
37     nspp = PossibleArrival - SimClasses.Clock
38     return nspp
39
40 def Arrival():
41     global MaxQueue
42     global N
43     interarrival = NSPP(1)
44     SimFunctions.Schedule(Calendar, "Arrival", interarrival)
45     N = N + 1
46     if N > MaxQueue:

```

```

47         MaxQueue = N
48         SimFunctions.Schedule(Calendar, "EndOfService", SimRNG.Expon(MeanBusyTime,
49                               2))
49
50     def EndOfService():
51         global N
52         N = N - 1
53
54     for Reps in range(0, RepNum, 1):
55         N = 0
56         MaxQueue = 0
57
58         SimFunctions.SimFunctionsInit(Calendar, TheQueues, TheCTStats, TheDTStats,
59                                       TheResources)
60         interarrival = NSPP(1)
61         SimFunctions.Schedule(Calendar, "Arrival", interarrival)
62         SimFunctions.Schedule(Calendar, "EndSimulation", 24*60)
63
64         NextEvent = Calendar.Remove()
65         SimClasses.Clock = NextEvent.EventTime
66         if NextEvent.EventType == "Arrival":
67             Arrival()
68         elif NextEvent.EventType == "EndOfService":
69             EndOfService()
70
71         while NextEvent.EventType != "EndSimulation":
72             NextEvent = Calendar.Remove()
73             SimClasses.Clock = NextEvent.EventTime
74             if NextEvent.EventType == "Arrival":
75                 Arrival()
76             elif NextEvent.EventType == "EndOfService":
77                 EndOfService()
78
79         AllMaxQueue.append(MaxQueue)
80
81     # estimating the 0.98th quantile
82     print("Estimated_required_drones_is:", np.sort(AllMaxQueue)[980])
83     l = int(np.floor(980 - 1.96*np.sqrt(980*0.02)))
84     u = int(np.ceil(980 + 1.96*np.sqrt(980*0.02)))
85     print("The_95%_CI: [{}, {}]".format(np.sort(AllMaxQueue)[l], np.sort(
86         AllMaxQueue)[u]))

```

Estimated required drones is: 224
The 95% CI: [222, 226]

Problem 3. (20 Pts.) The attached file ([TeslaPrices.csv](#)) provides daily prices for Tesla's stock during 2021.

(a) Fit a Geometric Brownian Motion (GBM) to the **Close** price by estimating the drift μ and volatility terms σ . Recall that for GBM the log-returns are independent Normal random variables, i.e.,

$$\log\left(\frac{S(t_i)}{S(t_{i-1})}\right) \sim N(\mu(t_i - t_{i-1}), \sigma(t_i - t_{i-1})),$$

for all $t_i, i = 1, \dots, k$.

We first calculate the log-returns using the **Close** price data. We then fit a normal distribution to the log-returns data and estimate $\mu = 0.001525905$ and $\sigma = 0.034080501$.

```
1 # This file uses the fitdistrplus package. You'll need to install it first
2 # by running the command: 'install.packages("fitdistrplus")' in R
3 require(fitdistrplus)
4
5 # import the data from .csv file
6 Data <- read.csv("TeslaPrices.csv", header=TRUE,
7                 sep=";", na.strings=c("NA", ""), stringsAsFactors=FALSE, as
8                 .is=TRUE)
9
10 Data$LogReturn = 0
11 for (i in 2:nrow(Data)){
12   Data$LogReturn[i] = log(Data$Close[i]/Data$Close[i-1])
13 }
14
15 MyData <- Data$LogReturn
16
17 # fit a normal distribution to the log-returns
18 fln <- fitdist(MyData, "norm")
19 fln$estimate
20
21 # Graphical method
22 par(mfrow = c(1, 2))
23 plot.legend <- c("normal")
24 denscomp(fln, legendtext = plot.legend)
25 qqcomp(fln, legendtext = plot.legend)
26
27 # Perform GofFit tests
28 gof_results <- gofstat(fln)
29 gof_results
30 gof_results$kstest
31 gof_results$chisq
32 gof_results$chisqpvalue
33 gof_results$chisqtable
```

mean	sd
0.001525905	0.034080501

(b) Does the GBM fit the data well? Use your estimates from part (a) together with both graphical methods and statistical goodness of fit tests discussed in class to justify your answer.

Yes, for this data set the Normal distribution seems appropriate. As shown in Fig. 5, both the Q-Q plot and histogram suggest that the Normal distribution is a good fit for the log-returns. Note that having larger deviation around the tails on the Q-Q plot is expected since the variance is higher at extreme quantiles. Also, the KS test does not reject the Normal distribution at $(1 - \alpha) = 95\%$ confidence level.

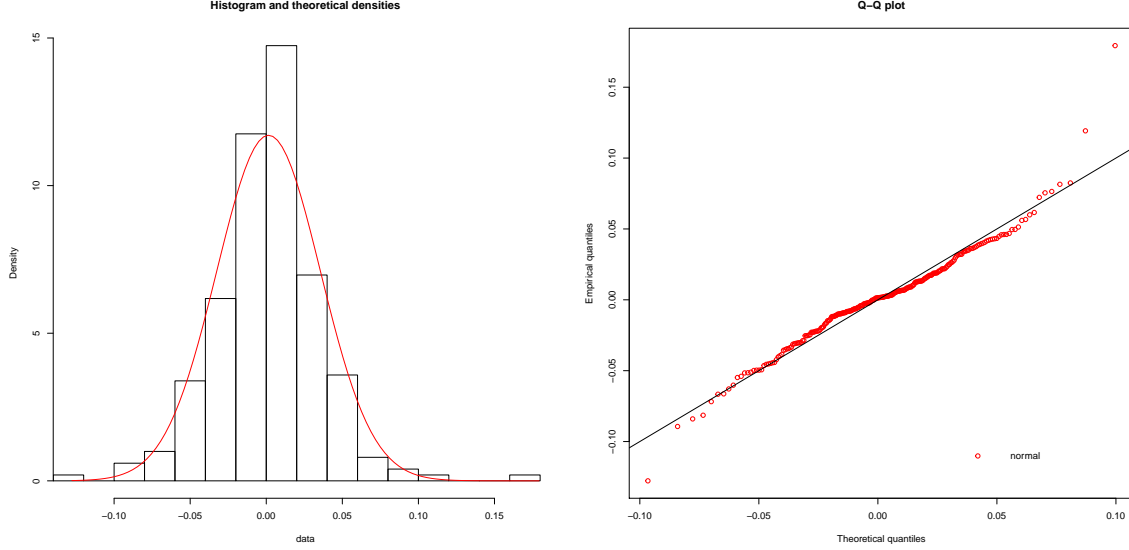


Figure 5: Graphical support for the goodness of fit of GBM.

Problem 4. (10 Pts.) Normal distribution is typically not a good fit for log-returns. Cauchy distribution is another (heavy-tailed) distribution used to model the log-returns. The CDF of the Cauchy distribution is given by:

$$F(x) = \frac{1}{\pi} \arctan\left(\frac{x - x_0}{\gamma}\right) + \frac{1}{2},$$

where x_0 and γ are parameters.

(a) Propose an inversion algorithm to generate samples of a Cauchy distribution with $x_0 = 0$ and $\gamma = 2$.

We denote the inverse cdf by $F^{-1}(\cdot)$. So,

$$F(x) = u \iff x = F^{-1}(u).$$

We have

$$\frac{1}{\pi} \arctan\left(\frac{x}{2}\right) + \frac{1}{2} = u \implies \arctan\left(\frac{x}{2}\right) = \pi\left(u - \frac{1}{2}\right) \implies x = 2 \tan\left(\pi\left(u - \frac{1}{2}\right)\right)$$

Therefore,

$$F^{-1}(u) = 2 \tan\left(\pi\left(u - \frac{1}{2}\right)\right)$$

So, the inversion algorithm works as follows:

1. Generate $U \sim U[0, 1]$.
2. Set $X = 2 \tan\left(\pi\left(U - \frac{1}{2}\right)\right)$ and return X .

(b) Estimating the parameters of the Cauchy distribution using MLE is challenging. One approach to estimate the parameters using matching is to match median (1/2th quantile) and half of the

inter-quartile range (difference between 1/4th and 3/4th quantiles) with their estimates from the data. Use this method to estimate the parameters using 1000 samples generated from part (a). How do they compare with the original parameters, i.e., $x_0 = 0$ and $\gamma = 2$?

We generate 1000 samples of the Cauchy distribution with $x_0 = 0$ and $\gamma = 2$ and estimate median and half of the inter-quartile range to be 0.05 and 3.98, respectively.

For the Cauchy distribution, we can calculate the q -th quantile using $x_0 + \gamma \tan(\pi(q - 1/2))$. So, matching median and half of the inter-quartile range with their estimates from the data, we have

$$x_0 + \gamma \tan(0) = 0.05 \implies x_0 = 0.05,$$

$$\gamma \tan\left(\frac{\pi}{4}\right) - \gamma \tan\left(-\frac{\pi}{4}\right) = 3.98 \implies \gamma = 1.99.$$

We observe that the estimated parameters are close enough to the original parameters.

```

1  # Cauchy distribution
2  import numpy as np
3
4  n = 1000
5  X_list = []
6  np.random.seed(1)
7  for i in range(n):
8      U = np.random.random()
9      X = 2*np.tan(np.pi*(U-1/2))
10     X_list.append(X)
11
12 print("Estimate_of_median:", np.median(X_list))
13 print("Estimate_of_difference_between_1/4th_and_3/4th_quantiles:", np.quantile
      (X_list, 3/4)-np.quantile(X_list, 1/4))

```

Estimate of median: 0.04713779707427136

Estimate of difference between 1/4th and 3/4th quantiles: 3.9842903109340155

Problem 5. (20 Pts.) The inversion method for generating a nonstationary arrival process requires an equilibrium base process with rate $\tilde{\lambda} = 1$ and variance σ_A^2 . When $\sigma_A^2 > 1$, one approach is to use a balanced hyperexponential distribution. This means that \tilde{A} is exponentially distributed with rate λ_1 with probability p , and exponentially distributed with rate λ_2 with probability $1 - p$. “balance” means that $p/\lambda_1 = (1 - p)/\lambda_2$. Thus, there are only two free parameters, p and λ_1 . The following values of p and λ can be shown to achieve the desired arrival rate and variance (Exercise 34, Chapter 6):

$$p = \frac{1}{2} \left(1 + \sqrt{\frac{\sigma_A^2 - 1}{\sigma_A^2 + 1}} \right), \lambda_1 = 2p.$$

(a) Propose an inversion method to generate arrivals according to the cumulative arrival rate $\Lambda(t) = t^2$ and with a balanced hyperexponential base with $\sigma_A^2 = 1.4$.

HINT: Start by deriving the distribution of the first inter-arrival time distribution $G_e(t)$ for the equilibrium renewal process with balanced hyperexponential base. The CDF of the hyperexponential distribution with parameters p, λ_1, λ_2 is given by,

$$G(t) = p(1 - e^{-\lambda_1 t}) + (1 - p)(1 - e^{-\lambda_2 t}).$$

We have $1 - G(t) = pe^{-\lambda_1 t} + (1 - p)e^{-\lambda_2 t}$, and therefore,

$$\begin{aligned} G_e(t) &= \tilde{\lambda} \int_0^t (1 - G(s)) ds \\ &= \int_0^t pe^{-\lambda_1 s} ds + \int_0^t (1 - p)e^{-\lambda_2 s} ds \\ &= \frac{p}{\lambda_1}(1 - e^{-\lambda_1 t}) + \frac{1 - p}{\lambda_2}(1 - e^{-\lambda_2 t}) \\ &= \frac{1}{2}(1 - e^{-\lambda_1 t}) + \frac{1}{2}(1 - e^{-\lambda_2 t}), \end{aligned}$$

where we have used $\tilde{\lambda} = 1$ and $p/\lambda_1 = (1 - p)/\lambda_2$. Note that the above implies that the initial inter-arrival time is equally likely to be exponential with rate λ_1 or λ_2 . Therefore, it has a hyperexponential distribution with $p = (1 - p) = 1/2$.

To generate a sample from a hyperexponential distribution, one can directly use the definition. In the general case, consider a hyperexponential distribution with parameters p, λ_1, λ_2 . That is, the random variable is exponentially distributed with rate λ_1 , with probability p , and exponentially distributed with rate λ_2 , with probability $(1 - p)$. Then one can generate a sample X as follows:

1. Generate $U_1 \sim Unif[0, 1]$ and $U_2 \sim Unif[0, 1]$
2. If $U_1 \leq p$: Return $X = \frac{-1}{\lambda_1} \ln(U_2)$; else: Return $X = \frac{-1}{\lambda_2} \ln(U_2)$.

Note that the above returns an exponential with rate λ_1 with probability p , and an exponential with rate λ_2 with probability $(1 - p)$. So, we generate the inter-arrival times A_2, A_3, \dots from the balanced hyper-exponential specified in the question and with parameters:

$$\begin{aligned} p &= \frac{1}{2} \left(1 + \sqrt{\frac{1.4 - 1}{1.4 + 1}} \right) = 0.7, \\ \lambda_1 &= 2p = 1.4, \\ \lambda_2 &= (1 - p)\lambda_1/p = 2(1 - p) = 0.6. \end{aligned}$$

The first inter-arrival time A_1 is generated from a hyper-exponential with

$$\begin{aligned} p &= 0.5, \\ \lambda_1 &= 1.4, \\ \lambda_2 &= 0.6. \end{aligned}$$

(b) Estimate $E(N(10))$ and $\text{Var}(N(10))$ using 100 samples generated using your algorithm in part (a) and compare the ratio $\text{Var}(N(10))/E(N(10))$ with the variance of the base equilibrium process. Explain your observation in a couple of sentences.

We generate 100 samples of $N(10)$ using the algorithm presented in part (a). The estimates of $E(N(10))$ and $\text{Var}(N(10))$ are 100.18 and 166.79, respectively. So, we estimate $\text{Var}(N(10))/E(N(10))$ to be 1.66 which is approximately close to the variance of the base equilibrium process.

```
1 import numpy as np
2 import matplotlib.pyplot as plt
3
```

```

4  # the function returns a random variate
5  # from hyperexponential(q,l1,l2)
6
7  def HypExp(q,l1,l2):
8      U1 = np.random.random()
9      U2 = np.random.random()
10     if U1 < q:
11         return (-1/l1)*np.log(U2)
12     else:
13         return (-1/l2)*np.log(U2)
14
15     np.random.seed(1)
16
17 # set parameters of the renewal process
18     p = 0.7
19     lambda1 = 2*p
20     lambda2 = 2*(1-p)
21
22     N10_list = []
23     for i in range(100):
24         # S-tilde denotes arrival times of the unit rate process
25         # generate the first inter-arrival from Ge(t)
26         S_tilde = HypExp(1/2, lambda1, lambda2)
27         # S denotes the actual arrival times
28         S = np.sqrt(S_tilde)
29         # N denotes the arrival-counting process
30         N = 0
31         # generate arrivals during [0,10]
32         while S <= 10:
33             N += 1
34             # generate the next arrival using G(t)
35             S_tilde = S_tilde + HypExp(p,lambda1,lambda2)
36             S = np.sqrt(S_tilde)
37
38         N10_list.append(N)
39
40     print("Estimate_of_the_Expectation:", np.mean(N10_list))
41     print("Estimate_of_the_Variance:", np.var(N10_list))
42     print("Estimate_of_the_Ration:", np.var(N10_list)/np.mean(N10_list))

```

```

Estimate of the Expectation: 100.18
Estimate of the Variance: 166.7876
Estimate of the Ration: 1.6648792174086642

```